

A FIELD GUIDE FOR EDITORS, REVIEWERS AND ANYONE WHO READS FOR A LIVING

The Tells

How machine-written prose gives itself away: the words, the sentence shapes and the habits of tone that language models fall into, why they fall into them, and a method for catching them on the page.

5 families of tells 100+ flagged words and phrases A worked example at the end

Why the machine writes like this

A language model predicts the next word. Each stage of its training pushes those predictions toward a narrow band of safe, pleasing, high-probability choices. The tells are what that narrowing leaves behind.

Stage 1	Stage 2	Stage 3	Stage 4
Pretraining The model reads a large slice of the internet. Marketing copy, SEO blogs, press releases, LinkedIn posts and corporate reports are heavily represented, so their vocabulary becomes the model's idea of normal. Leaves: robust, seamless, landscape, leverage, in today's fast-paced world	Instruction tuning The model learns what a "good answer" looks like from example dialogues. Those examples favor headers, bullet lists, a summary at the end and an offer to help further. Leaves: bold-lead bullets, "In summary," "Let me know if..."	Human feedback Raters score pairs of answers. Answers that sound confident, warm, thorough and insightful win. The model learns which flourishes earn points, whether or not they add information. Leaves: "Great question!", "not X but Y", inflated significance	Sampling When generating, the model favors high-probability words. The safe, middle-of-the-road phrasing wins over and over, which makes every output sound like the average of every output. Leaves: the same 200 words, the same rhythm, every time

↳ Feedback loop: AI text gets published, ends up in the next model's training data, and creeps into people's own writing. Each cycle concentrates the habits further.

Two ideas explain almost every tell in this guide.

Regression to the mean. A model picks the most expected word. Human writers pick the right word for this sentence, which is often unexpected. Text made entirely of expected choices has a recognizable flatness, even when every sentence is grammatical.

Performing helpfulness. The model was rewarded for sounding insightful, so it learned the *shapes* of insight: the reversal, the triplet, the dramatic short sentence, the claim that something matters deeply. These shapes appear whether or not there is an insight to carry.

Most AI tells are the packaging of an idea delivered without the idea.

Sentence shapes

These are the strongest tells because they survive word substitution. Someone can swap "delve" for "explore," but the skeleton of the sentence stays the same.

THE TELL	SPECIMEN	WHY THE MODEL DOES IT	PLAIN VERSION
Not X, but Y (the false reversal)	"It's not just a survey. It's a conversation with your community." "This isn't about data. It's about people."	It invents a position nobody holds, then corrects it. The correction feels like a revelation, though no new information was given. Raters reward the feeling.	<i>State the claim once: "The survey asks residents what they want."</i>
The colon or question reveal	"The result? Faster decisions." "Here's the thing: trust is earned." "The kicker:"	A drumroll before an ordinary statement. It copies the cadence of a keynote speaker and borrows that authority.	<i>"Decisions came faster."</i>
Rule of three	"Clear, concise, and compelling." "Faster, smarter, more efficient."	Triplets sound complete and rhythmic. Models produce them by reflex, even when only one or two of the items are true.	<i>Keep the item that is true and specific.</i>
The dramatic fragment	"...and the council voted it down. Unanimously." "And that changes everything."	Copied from viral social posts. The short line after a long one signals emphasis without earning it.	<i>Fold it into the sentence, or cut it.</i>

THE TELL	SPECIMEN	WHY THE MODEL DOES IT	PLAIN VERSION
The tacked-on "-ing" clause	"The city launched the program, highlighting its commitment to transparency." ", underscoring the importance of..."	A cheap way to add significance to a plain fact. The participle asserts meaning the writer never established.	<i>"The city launched the program." Stop there.</i>
Copula dodging	"serves as a hub" · "stands as a reminder" · "boasts 40 parks" · "represents a shift"	Models avoid plain "is" and "has," likely because varied verbs scored as more sophisticated.	<i>"is a hub" · "has 40 parks"</i>
The false range	"From boardrooms to classrooms..." "Whether you're a startup or a Fortune 500..."	Sounds sweeping and inclusive. The two endpoints rarely define a real spectrum.	<i>Name who it is actually for.</i>
Signposting	"Let's dive in." · "Let's break it down." · "Here's what you need to know."	Narrates the structure of the text instead of delivering the content.	<i>Delete. Start with the first real point.</i>
Synonym cycling	"the city... the municipality... the local government... the jurisdiction"	A repetition penalty pushes the model to vary nouns. People reuse the same word for the same thing.	<i>Pick one name and keep it.</i>
The summary ending	"Ultimately, ..." · "In short, ..." · "At the end of the day, ..." · "In conclusion, ..."	Instruction-tuning examples end with a recap. The model recaps even three-paragraph answers.	<i>End on the last new point.</i>

The lexicon

Words that appear in AI text far more often than in human writing on the same subject. One of them proves nothing. Five in a paragraph is a pattern. Each word is followed by a plain substitute.

VERBS

delve *look at*

leverage *use*

utilize *use*

landed *worked, was received*

surfaced *found, raised*

navigate *handle*

foster *build, encourage*

underscore *show*

showcase *show*

streamline *simplify*

elevate *improve*

unlock *allow, open*

unpack *explain*

harness *use*

empower *let, help*

resonate *matter to*

embark *start*

bolster *support*

NOUNS

landscape *field, market*

realm *area*

tapestry *cut it*

testament *proof*

journey *process*

ecosystem *name the parts*

cornerstone *basis*

interplay *relationship*

insights *findings*

synergy *cut it*

paradigm *model*

game- *say what*

changer *changed*

north star *goal*

beacon *cut it*

ADJECTIVES

robust *strong, or say how*

seamless *say what didn't break*

pivotal *important, or cut*

crucial / vital *needed*

comprehensive *complete*

intricate *complicated*

meticulous *careful*

nuanced *state the nuance*

dynamic / vibrant *cut it*

transformative *say what changed*

holistic *whole*

actionable *usable*

compelling *convincing*

INTENSIFIERS & FILLER		NEWER TELLS (2025-26)		STOCK PHRASES	
really	<i>cut</i>	load-bearing	<i>essential</i>	in today's fast-paced world	<i>cut</i>
simply	<i>cut</i>	lands / landed	<i>works</i>	ever-evolving	<i>changing</i>
truly	<i>cut</i>	surfaces	<i>shows</i>	plays a pivotal role	<i>matters because...</i>
genuinely	<i>cut</i>	the real question	<i>the question</i>	a significant milestone	<i>name it</i>
honestly	<i>cut</i>	doing a lot of work	<i>say what it does</i>	stands as a testament	<i>shows</i>
straightforward	<i>simple</i>	lean into	<i>commit to</i>	key takeaways	<i>findings</i>
incredibly	<i>cut</i>	sharp / crisp	<i>clear</i>	at the end of the day	<i>cut</i>
deeply	<i>cut</i>	earned	<i>justified</i>	moving forward	<i>next</i>
quietly	<i>cut</i>	in practice	<i>cut</i>	a deep dive	<i>a review</i>
fundamentally	<i>cut</i>	the kicker	<i>cut</i>		
notably	<i>cut</i>				
arguably	<i>make the argument</i>				
it's worth noting	<i>cut</i>				
it's important to note	<i>cut</i>				

Why "landed" and "surfaced"? Newer models are trained heavily on software-team writing: commit messages, Slack threads, product reviews. That jargon ("the change landed," "the tool surfaced an error") leaks into ordinary prose. The intensifiers ("really," "simply," "genuinely") come from a different source: the model reaching for sincerity and emphasis it can't show through content.

Tone and stance

The voice of a model trained to please: agreeable, upbeat, careful to offend no one, and convinced everything is significant.

THE TELL	SPECIMEN	WHAT'S GOING ON	PLAIN VERSION
Sycophantic opener	"Great question!" • "Absolutely!" • "You're right to think about this."	Praise scored well with raters. It now appears before any answer, regardless of the question.	<i>Answer the question.</i>
Service closer	"I hope this helps!" • "Let me know if you'd like me to..." • "Feel free to reach out."	Chat-assistant residue. It often survives copy-paste into emails and reports.	<i>Delete.</i>
Inflated significance	"marks a pivotal moment" • "a powerful reminder of" • "reshaping how we think about"	Everything is framed as historic because importance reads as insight.	<i>Give the fact and let the reader judge it.</i>
Hedge stacking	"This could potentially help to some extent in certain situations."	Trained to avoid being wrong, the model qualifies every claim.	<i>One qualifier, where it's earned.</i>
Reflexive balance	"While there are benefits, there are also challenges to consider."	Both-sides framing on questions that have a clear answer.	<i>Take the position the evidence supports.</i>
Vague authority	"Experts agree..." • "Studies show..." • "Many believe..."	The shape of a citation with no source behind it.	<i>Name the expert or study, or drop the claim.</i>
Moral coda	"...reminding us that, in the end, listening is the most powerful tool we have."	An uplifting closer added because endings "should" resonate.	<i>End on the last fact.</i>

THE TELL	SPECIMEN	WHAT'S GOING ON	PLAIN VERSION
Promotional sheen	"nestled in the heart of" • "rich cultural heritage" • "a must-visit"	Travel-brochure and press-release diction absorbed in pretraining.	<i>Describe what is there.</i>

PART 2 • FAMILY 4

Formatting habits

Chat interfaces render Markdown, so models learned to structure everything. The structure follows them into places it doesn't belong.

THE TELL	SPECIMEN	WHAT'S GOING ON
Bold-lead bullets	• Speed: Faster results. • Clarity: Clearer data. • Trust: Better outcomes.	Every bullet the same length and shape, each starting with a bolded label and a colon. Real notes are uneven.
Headers on short answers	## Overview ## Key Benefits ## Conclusion	Three paragraphs carrying three headers, in Title Case, often ending in a "Conclusion" or "Challenges and Future Outlook" section.
Em-dash density	"The data – collected over six weeks – shows a clear trend – one that matters."	People use em dashes too. Three or four per paragraph, unspaced, in someone who never used them before, is the signal.
Emoji markers	 Launch  Benefits  Tip	Common in LinkedIn-style output and in lists generated for social posts.
Uniform rhythm	Paragraph lengths: 4, 4, 4, 4 sentences. Sentence lengths: 18-24 words, every one.	Researchers call this low "burstiness." Human writing swings between short and long.
Markdown leaking	**important** in a plain-text email "### Next steps" in a Word doc	Raw formatting characters pasted from a chat window.

Smoking guns

Most tells are probabilistic. These are near-certain: residue of the chat session that someone forgot to remove.

ARTIFACT	SPECIMEN	WHERE IT COMES FROM
Self-reference	"As an AI language model..." • "As of my last knowledge update..." • "I don't have access to real-time data."	The model describing its own limits.
Answer preamble	"Certainly! Here's a revised version of your paragraph:"	The chat reply copied along with the content.
Unfilled placeholders	[Your Name] • [Company] • [Insert statistic here]	Template slots the user never filled in.
Citation debris	utm_source=chatgpt.com • oaicite • turn0search3 • [4↑source]	Tracking parameters and internal reference markers from AI search tools.
Fake or wrong citations	A real journal, a real-sounding author, a DOI that resolves to nothing.	The model producing the shape of a reference. Check any citation you can't verify.

How to spot it: a six-step read

No single word convicts a text. Read for density and for the absence of specifics. These steps run in order, from fastest to slowest.

1 Check the edges

Read the first and last sentences. Praise openers, recap closers, uplifting codas and "let me know" offers cluster there.

2 Count the clusters

Circle every lexicon word and sentence shape. One per page is normal writing. Three or more per paragraph is a pattern.

3 Run the swap test

Could this sentence appear unchanged in a document about a different city, product or topic? If yes, it carries no information about this one.

4 Hunt for specifics

Look for names, dates, figures, places, quotes and first-hand detail. AI text is long on claims and short on particulars that could be checked.

5 Read it aloud

AI prose has an even, sing-song rhythm: matched sentence lengths, triplets, and a short punchline after each long sentence. Your ear catches this faster than your eye.

6 Compare to the writer

The best evidence is a change. If someone who writes in plain, short emails suddenly sends four tidy paragraphs with em dashes, that shift says more than any word list.

Worked example

Below is an example paragraph written to show the tells, run through a simple scanner built on the word lists and sentence patterns in this guide. Solid blocks are flagged words and phrases; underlines are sentence shapes.

WORDS	TELLS	PER 100 WORDS
119	29	24.4

VERY HIGH: SATURATED WITH TELLs

word lexicon or phrase **shape** sentence structure or artifact

Great question! **In today's fast-paced** **landscape**, local governments aren't just collecting data — they're **unlocking** **insights** that **truly** **resonate** with residents.

Here's the thing: a **robust** survey program **serves as** a **cornerstone** of public trust, fostering transparency and **empowering** decision-makers to **navigate** complex challenges. The result? Findings that **landed** with council members and **surfaced** priorities that might otherwise go unnoticed.

It's worth noting that this isn't about numbers. It's about people. **Whether you're** a small town or a major county, a **comprehensive** approach can be a **game-changer** for your community's **journey**.

Ultimately, polling **stands as a testament** to the power of listening. **I hope this helps!** **Let me know if** you'd like me to expand on any of these points.

Found: reveal ×2 · great question ×1 · in today's fast-paced ×1 · landscape ×1 · not just ×1 · em dash ×1 · unlocking ×1 · insights ×1 · truly ×1 · resonate ×1 · robust ×1 · serves as ×1 · cornerstone ×1 · -ing tack-on ×1 · empowering ×1 · navigate ×1 · landed ×1 · surfaced ×1 · it's worth noting ×1 · not X. It's Y. ×1 · whether you're ×1 · comprehensive ×1 · game-changer ×1 · journey ×1 · ultimately ×1 · stands as a testament ×1 · i hope this helps ×1 · let me know if ×1

A density above about 3 tells per 100 words is a strong pattern. Word lists miss tells they have no pattern for and will flag human writers who use these words honestly, so use a count like this to direct your attention, then judge with the six-step read.

Before you accuse anyone

People wrote these words first

Models learned "robust," the em dash and the rule of three from human writers. Consultants, academics and marketers used them for decades. A tell is evidence of a habit, and the habit predates AI.

Editing hides the surface

A careful person can strip the lexicon from AI text in minutes. Sentence shapes and the lack of specifics are harder to remove, which is why steps 3 and 4 matter most.

The useful question

Whether a machine wrote it matters less than whether the text says something specific, true and checkable. Every tell in this guide is also a sign of weak writing, whoever produced it.

Detectors are unreliable

Automated AI detectors produce false positives often enough that several universities have turned them off. A 2023 Stanford study found they flagged most essays by non-native English speakers as AI-written.

The habits are spreading

Researchers found "delve" and similar words surging in scientific abstracts after 2023, at rates that suggest at least one in ten papers was processed with an LLM. People who read a lot of AI text start to write like it.

SOURCES AND FURTHER READING

1. Kobak, González-Márquez, Horvát & Lause, "Delving into LLM-assisted writing in biomedical publications through excess vocabulary," *Science Advances* (2025).
2. Liang et al., "GPT detectors are biased against non-native English writers," *Patterns* (2023).
3. Wikipedia, "[Signs of AI writing](#)," maintained by WikiProject AI Cleanup.